

Disaster Tweets

Applying Natural Language Processing (NLP) techniques



Gaurav Dembla

Nov 7, 2020



Problem Statement

Predict which Tweets are about real disasters and which ones aren't

- Twitter, social networking service, allows users to post messages known as “tweets”
- Users observe or experience a disaster -> post tweets in real-time
- Person's words clear to a human right away, but not so much to a machine
- Example: “On plus side LOOK AT THE SKY LAST NIGHT IT WAS ABLAZE”.
Here, the author explicitly uses the word “ABLAZE”, but means it metaphorically.

Source Data

- **Source:** Kaggle
- **Size:** 7613 x 5
- *Target* value of 1 denotes tweet as a disaster
- **Success Criteria:** ROC-AUC (Area Under the Curve of Receiver Operating Characteristic graph) on unseen data (hold-out/test set)

id	keyword	location	text	target
1			Our Deeds are the Reason of this #earthquake May ALLAH Forgive us all	1
4			Forest fire near La Ronge Sask. Canada	1
5			All residents asked to 'shelter in place' are being notified by officers. No other evacuation or shelter in place orders are expected	1
48	ablaze	Birmingham	@bbcmtd Wholesale Markets ablaze http://t.co/lHYXEOHY6C	1
49	ablaze	Est. September 2012 - Bristol	We always try to bring the heavy. #metal #RT http://t.co/YAo1e0xngw	0
50	ablaze	AFRICA	#AFRICANBAZE: Breaking news:Nigeria flag set ablaze in Aba. http://t.co/2nndBGwyEi	1
112	accident	San Mateo County, CA	Traffic accident N CABRILLO HWY/MAGELLAN AV MIR (08/06/15 11:03:58)	1
119	accident		Can wait to see how pissed Donnie is when I tell him I was in ANOTHER accident??	0

Exploratory Data Analysis

- **Data Balance:** Nearly balanced -> No data balancing required

Target	% of records
0	57.0%
1	43.0%

- **Missing Values**

- Location with 1/3 values -> Drop the attribute
- Keyword is redundant information -> Drop the attribute

Attribute	% Nulls
Keyword	0.8%
Location	33.2%
Text	0.0%

- **URLs (http://)** and **mentions (@someone)** are of no use whatsoever
- **Emoticons** could be useful in predictions, as humans tend to use them in casual language to express their feelings and state of mind

id	keyword	location	text	target
1			Our Deeds are the Reason of this #earthquake May ALLAH Forgive us all	1
4			Forest fire near La Ronge Sask. Canada	1
40			Cooooo! :)	0
48	ablaze	Birmingham	@bbcmdt Wholesale Markets ablaze http://t.co/lHYXEOHY6C	1
49	ablaze	Est. September 2012 - Bristol	We always try to bring the heavy. #metal #RT http://t.co/YAo1e0xngw	0
50	ablaze	AFRICA	#AFRICANBAZE: Breaking news:Nigeria flag set ablaze in Aba. http://t.co/2nndBGwyEi	1
112	accident	San Mateo County, CA	Traffic accident N CABRILLO HWY/MAGELLAN AV MIR (08/06/15 11:03:58)	1
119	accident		Can wait to see how pissed Donnie is when I tell him I was in ANOTHER accident??	0

Data Cleaning

Remove the following items:

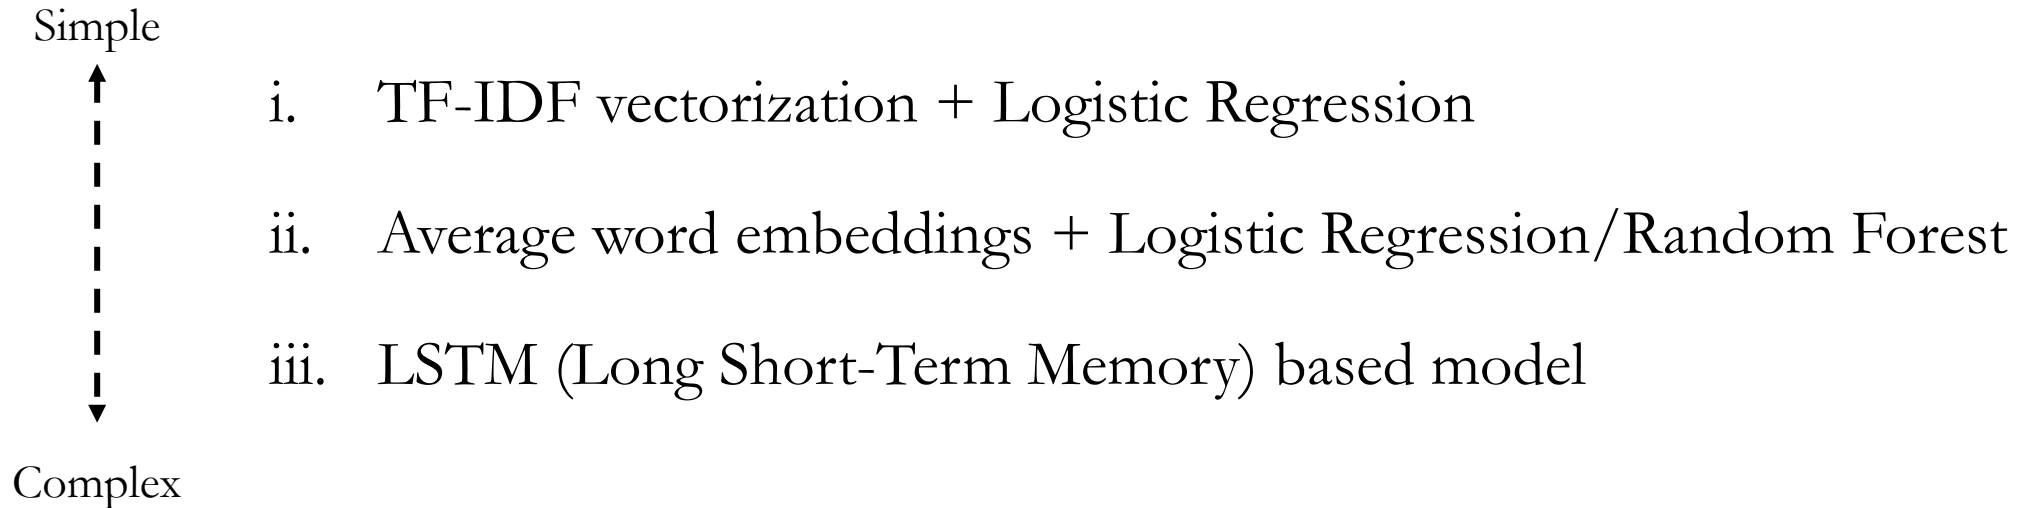
- ✓ URLs
- ✓ Mentions (@)
- ✓ Emoticons
- ✓ Punctuations (! # \$ % & ' () * ^ _ - + / \ < = > [] { } | ~ . “ ` ‘ , : ;). Retain hashtag text (string following # character)
- ✓ Tabs and line breaks
- ✓ Numeric digits
- ✓ Stop words using nltk library. E.g. “the”, “is”, “in”, “for”, “where”, “when”, “to”, “at” etc.
- ✓ Non-ascii characters. E.g. convert Carolinaâ€™Ablaze to CarolinaAblaze.

Apply following transformations:

- ✓ Lower case the characters - “Fire” and “fire”.
- ✓ Lemmatize using spacy library

Word	Lemma
seen/saw/seeing/see	see
drove/drive/driving	drive
better/good	good
playing/played/play	play

Vectorization/Statistical Modeling



Loss function: Log-loss (aka binary cross-entropy) to optimize the model

TF-IDF

Vectorization

Clean Documents

Pre-processing

Remove –

High-freq. words (present in more than 80% of documents)

Low-freq. words (present in less than 5 documents)

- d_1 : "sky blue"
- d_2 : "sun bright today"
- d_3 : "sun sky bright"
- d_4 : "can see shining sun bright sun"

$f_{t,d}$

	blue	bright	can	see	shining	sky	sun	today
1	1	0	0	0	0	1	0	0
2	0	1	0	0	0	0	1	1
3	0	1	0	0	0	1	1	0
4	0	1	1	1	1	0	2	0

$N=4$

$$\text{tf}(t, d) = \frac{f_{t,d}}{\sum_{t'} f_{t',d}}$$

	blue	bright	can	see	shining	sky	sun	today
1	1/2	0	0	0	0	1/2	0	0
2	0	1/3	0	0	0	0	1/3	1/3
3	0	1/3	0	0	0	1/3	1/3	0
4	0	1/6	1/6	1/6	1/6	0	1/3	0

$\times =$

$$\text{idf}(t, D) = \log_{10} \frac{N}{n_t}$$

$$\text{tfidf}(t, d, D) = \text{tf}(t, d) \cdot \text{idf}(t, D)$$

	blue	bright	can	see	shining	sky	sun	today
1	0.301	0	0	0	0	0.151	0	0
2	0	0.0417	0	0	0	0	0.0417	0.201
3	0	0.0417	0	0	0	0.100	0.0417	0
4	0	0.0209	0.100	0.100	0.100	0	0.0417	0

blue	bright	can	see	shining	sky	sun	today
0.602	0.125	0.602	0.602	0.602	0.301	0.125	0.602

$$\log_{10} \frac{4}{1} = 0.602$$

$$\log_{10} \frac{4}{3} = 0.125$$



TF-IDF

Statistical Modeling

$$\text{tfidf}(t, d, D) = \text{tf}(t, d) \cdot \text{idf}(t, D)$$

	blue	bright	can	see	shining	sky	sun	today
1	0.301	0	0	0	0	0.151	0	0
2	0	0.0417	0	0	0	0	0.0417	0.201
3	0	0.0417	0	0	0	0.100	0.0417	0
4	0	0.0209	0.100	0.100	0.100	0	0.0417	0

Sparse vector of 1783 length



Logistic
Regression

Average Word Embedding

Vectorization

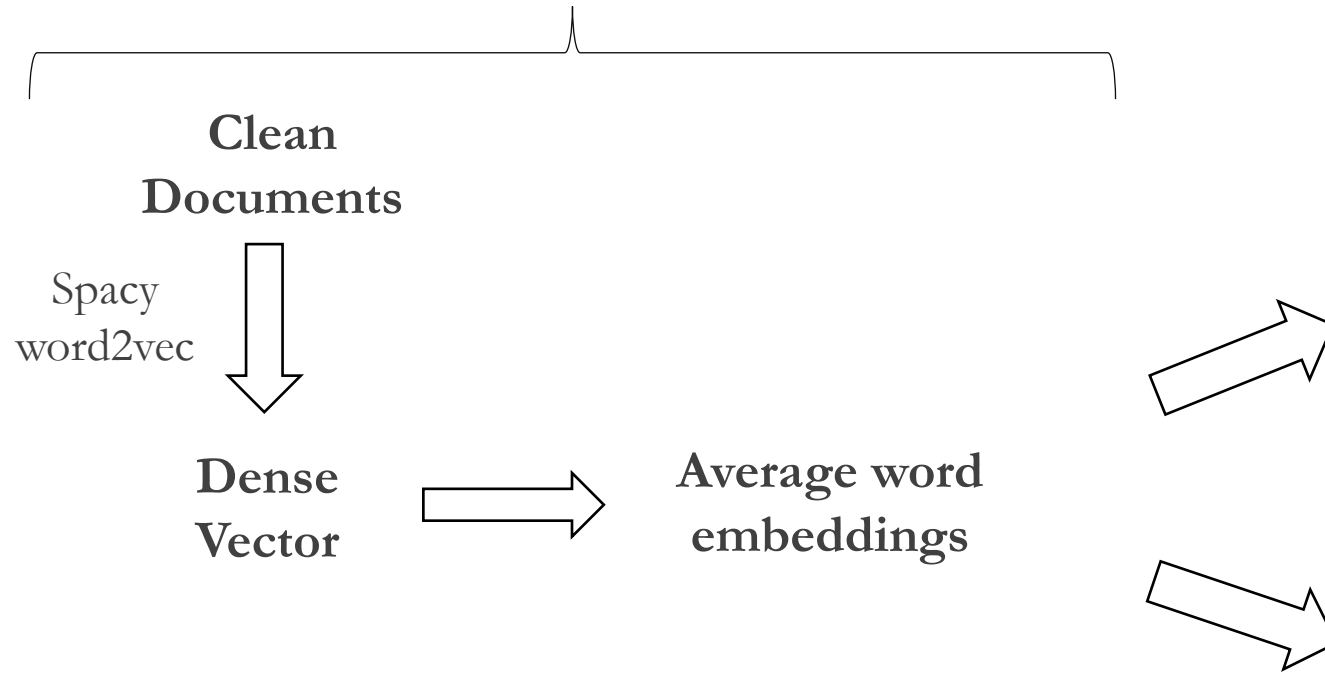


Illustration:

Document: “there is fire on the hill”



Remaining tokens after removing stop-words: [fire, hill]

List of word embeddings (spacy): [[3, 6, 2], [9, 4, 6]]

Average embedding of the document: $[(3+9)/2, (6+4)/2, (2+6)/2] = [6, 5, 4]$

Statistical Modeling

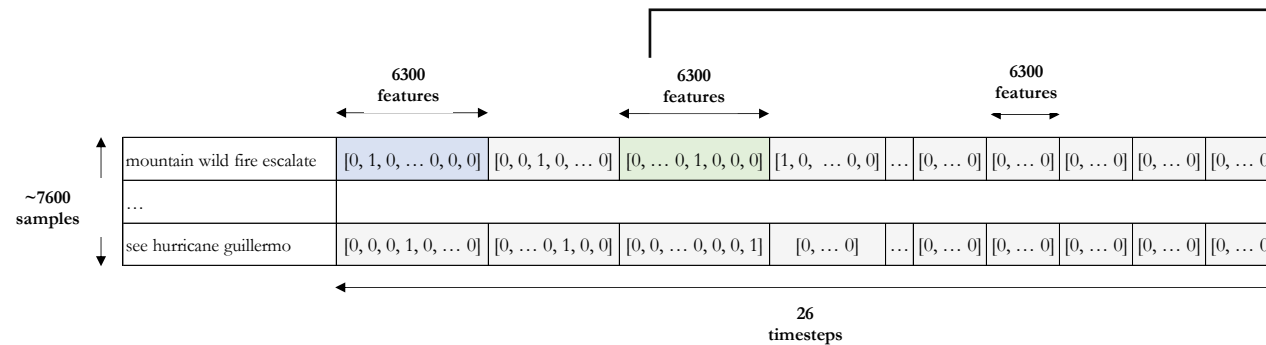
Logistic Regression



Random Forest

LSTM

Vectorization



Sparse vectorization

Document (clean text)	0	1	2	3	...	20	21	22	23	24	25
mountain wild fire escalate	406	2253	2254	1577	...	0	0	0	0	0	0

Fixed-length sequencing

Document (clean text)	Sequence of Indexes
mountain wild fire escalate	[406, 2253, 2254, 1577]

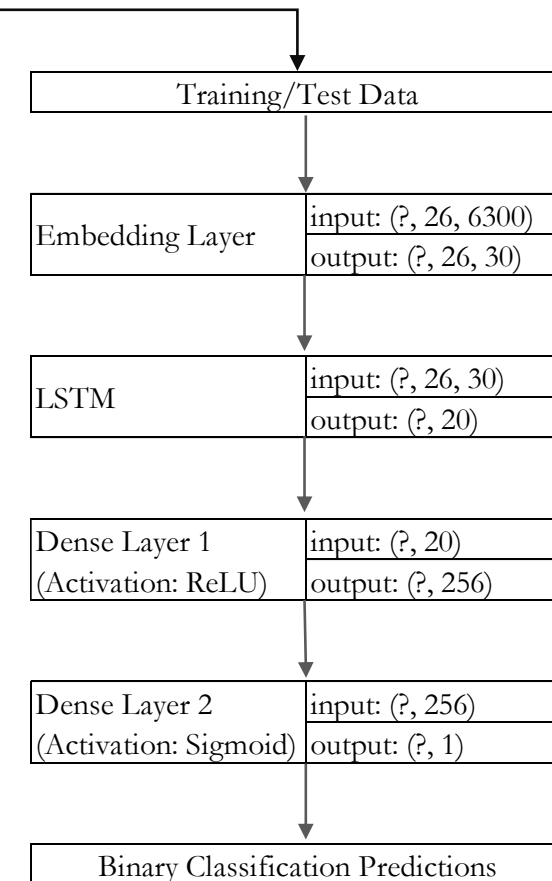
(Variable-length) Sequencing

$[('mountain', 406), ('wild', 2253), ('fire', 2254), ('escalate', 1577), ('sky', 213), ('region', 24)]$

Build vocabulary (word, index pairs)

Corpus (clean documents)


Statistical Modeling



Batch Size: 64

Epochs: 20

(with early-stopping criteria)



Conclusion

<i>Vectorization Model</i>	TF-IDF	Document-level Word Embedding		LSTM
<i>Statistical Model</i>	Logistic Regression	Logistic Regression	Random Forest	
ROC AUC	0.847	0.853	0.858	0.851
Log-loss	0.4685	0.4676	0.4713	0.4579
Accuracy	78.7%	79.1%	81.0%	79.6%
Training Effort (in minutes)	1	1	17	1

✓ Document-level word embedding + Random Forest

- Simple to comprehend
- Best results of the lot
 - Training effort large relative to others

Future Scope

Transfer learning: Existing LSTM-based model – Embedding layer + Spacy word2vec

[Project Repository](#)

[Detailed Report](#)

Code: [Disaster Tweets.ipynb](#) (Jupyter Notebook)